

Problem Statement

You are working for a new-age insurance company and employ multiple outreach plans to sell term insurance to your customers. Telephonic marketing campaigns still remain one of the most effective way to reach out to people however they incur a lot of cost. Hence, it is important to identify the customers that are most likely to convert beforehand so that they can be specifically targeted via call. We are given the historical marketing data of the insurance company and are required to build a ML model that will predict if a client will subscribe to the insurance.

Features:

age (numeric) job : type of job marital : marital status educational_qual : education status call_type : contact communication type day: last contact day of the month (numeric) mon: last contact month of year dur: last contact duration, in seconds (numeric) num_calls: number of contacts performed during this campaign and for this client prev_outcome: outcome of the previous marketing campaign (categorical: "unknown", "other", "failure", "success")

Output variable (desired target):

y - has the client subscribed to the insurance

Importing Libraries and Data

In [1]:

```
1  #Importing Libraries
2
3  import pandas as pd
4
5  import numpy as np
6
7  import seaborn as sns
8
9  import matplotlib.pyplot as plt
10 %matplotlib inline
11
12
13 # For ignore warnings
14 import warnings
15 warnings.filterwarnings("ignore")
```

In [2]:

```

1 #Importing data
2 data = pd.read_csv(r"D:\IITM_Final_Project\Customer Conversion Prediction - Customer
3 data.info()

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 45211 entries, 0 to 45210
```

```
Data columns (total 11 columns):
```

#	Column	Non-Null Count	Dtype
0	age	45211 non-null	int64
1	job	45211 non-null	object
2	marital	45211 non-null	object
3	education_qual	45211 non-null	object
4	call_type	45211 non-null	object
5	day	45211 non-null	int64
6	mon	45211 non-null	object
7	dur	45211 non-null	int64
8	num_calls	45211 non-null	int64
9	prev_outcome	45211 non-null	object
10	y	45211 non-null	object

```
dtypes: int64(4), object(7)
```

```
memory usage: 3.8+ MB
```

In [3]:

```
1 data.head()
```

Out[3]:

	age	job	marital	education_qual	call_type	day	mon	dur	num_calls	prev_out
0	58	management	married	tertiary	unknown	5	may	261	1	unl
1	44	technician	single	secondary	unknown	5	may	151	1	unl
2	33	entrepreneur	married	secondary	unknown	5	may	76	1	unl
3	47	blue-collar	married	unknown	unknown	5	may	92	1	unl
4	33	unknown	single	unknown	unknown	5	may	198	1	unl

In [4]:

```
1 data.columns
```

Out[4]:

```
Index(['age', 'job', 'marital', 'education_qual', 'call_type', 'day', 'mo
n',
      'dur', 'num_calls', 'prev_outcome', 'y'],
      dtype='object')
```

In [5]:

```
1 data.shape
```

Out[5]:

```
(45211, 11)
```

In [6]:

```
1 data.dtypes
2 #object are categorical datatypes
```

Out[6]:

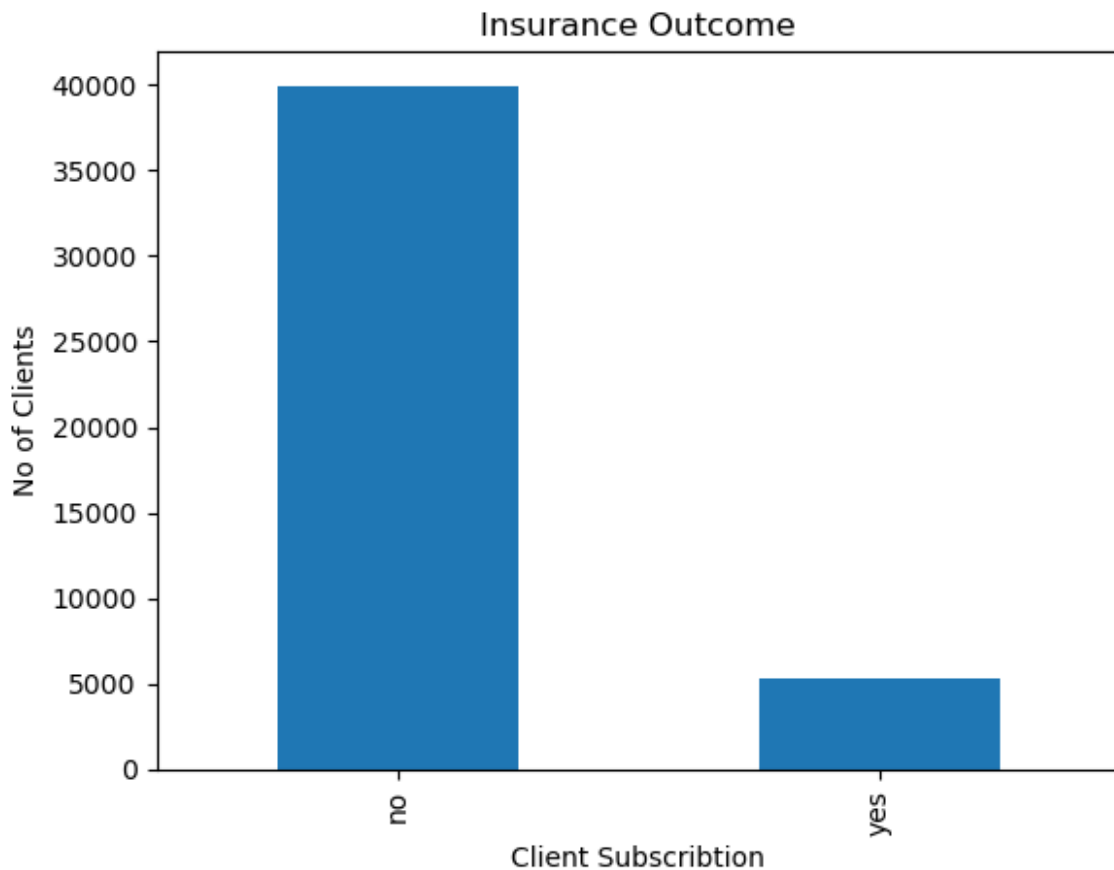
```
age          int64
job          object
marital      object
education_qual  object
call_type    object
day          int64
mon          object
dur          int64
num_calls    int64
prev_outcome object
y            object
dtype: object
```

In [7]:

```
1 #Target variable analysis
2 print(data['y'].value_counts())
3
4 print('\nNot Subscribed - ', round(data['y'].value_counts()['no']/len(data)*100, 2),
5 print('Subscribed - ', round(data['y'].value_counts()['yes']/len(data)*100, 2), '% c
6
7 #Plotting distribution
8 data['y'].value_counts().plot(kind='bar')
9 plt.title('Insurance Outcome')
10 plt.xlabel('Client Subscription')
11 plt.ylabel('No of Clients')
12 plt.show()
```

```
no      39922
yes      5289
Name: y, dtype: int64
```

Not Subscribed - 88.3 % of the dataset
Subscribed - 11.7 % of the dataset



Data Cleaning

In [8]:

```
1 # checking for null values
2 data.isnull().sum()
```

Out[8]:

```
age          0
job          0
marital      0
education_qual  0
call_type    0
day          0
mon          0
dur          0
num_calls    0
prev_outcome 0
y           0
dtype: int64
```

No null values in the data, but we can see there are 'unknown' category which we consider as null values.

In [9]:

```
1 data[data=='unknown'].count()
```

Out[9]:

```
age          0
job          288
marital      0
education_qual 1857
call_type    13020
day          0
mon          0
dur          0
num_calls    0
prev_outcome 36959
y           0
dtype: int64
```

In [10]:

```
1 # Dropping the duplicate values
2 data.drop_duplicates(inplace=True)
```

Job Column

In [11]:

```
1 data['job'].unique()
```

Out[11]:

```
array(['management', 'technician', 'entrepreneur', 'blue-collar',
      'unknown', 'retired', 'admin.', 'services', 'self-employed',
      'unemployed', 'housemaid', 'student'], dtype=object)
```

In [12]:

```
1 # Count of unknown values
2 job_uc = data['job'][data['job']=='unknown'].count()
3 print("unknown job",job_uc)
4 #percentage of unknown values
5 job_ucp = (job_uc/len(data['job']))*100
6 print(f'{round(job_ucp, 2)}% of job column is unknown')
```

unknown job 288

0.64% of job column is unknown

In [13]:

```
1 #Imputing Job Column
2 # Imputing job column
3 data['job'].replace('unknown', data['job'].mode()[0], inplace=True)
4 data['job'].unique()
```

Out[13]:

```
array(['management', 'technician', 'entrepreneur', 'blue-collar',
      'retired', 'admin.', 'services', 'self-employed', 'unemployed',
      'housemaid', 'student'], dtype=object)
```

Marital Column

In [14]:

```
1 data.marital.unique()
2 # no unknown values
```

Out[14]:

```
array(['married', 'single', 'divorced'], dtype=object)
```

education_qual column

In [15]:

```
1 data.education_qual.unique()
```

Out[15]:

```
array(['tertiary', 'secondary', 'unknown', 'primary'], dtype=object)
```

In [16]:

```
1 eq_uc = data['education_qual'][data['education_qual']=='unknown'].count()
2 print("unknown education_qual",eq_uc)
3 eq_ucp = eq_uc/len(data['education_qual'])*100
4 print(f'{round(eq_ucp, 2)}% of education_qual column is unknown')
5 # Imputing education_qual Column
6 data['education_qual'].replace('unknown', data['education_qual'].mode()[0], inplace=
```

unknown education_qual 1857

4.11% of education_qual column is unknown

Exploratory Data Analysis

In [17]:


```
1 num_col=list(data._get_numeric_data())
2 cat_col=list(data.select_dtypes(include = ['object']))
3 print("Numerical columns are",num_col)
4 print("Categorical columns are",cat_col)
```

Numerical columns are ['age', 'day', 'dur', 'num_calls']
Categorical columns are ['job', 'marital', 'education_qual', 'call_type', 'mon', 'prev_outcome', 'y']

In [18]:


```
1 for i in num_col:
2     plt.figure()
3     sns.displot(data=data[i],kde=True,bins=50)
```

<Figure size 640x480 with 0 Axes>



In [19]:

```
1 for i in cat_col:
2     plt.figure(figsize=(15,5)),
3     sns.countplot(x=data[i],data=data)
4
```



Encoding

In [20]:

```
1 #Target column encoding
2 data['y'] = data['y'].map({'yes':1, 'no':0})
3 data.head()
```

Out[20]:

	age	job	marital	education_qual	call_type	day	mon	dur	num_calls	prev_out
0	58	management	married	tertiary	unknown	5	may	261	1	unl
1	44	technician	single	secondary	unknown	5	may	151	1	unl
2	33	entrepreneur	married	secondary	unknown	5	may	76	1	unl
3	47	blue-collar	married	secondary	unknown	5	may	92	1	unl
4	33	blue-collar	single	secondary	unknown	5	may	198	1	unl

In [21]:

```
1 #Job column encoding
2 job_cat_ord = {'blue-collar':1, 'entrepreneur':2, 'housemaid':3, 'services':4, 'tech
3               'admin.':7, 'management':8, 'unemployed':9, 'retired':10, 'student':11
4 data['job'] = data['job'].map(job_cat_ord)
5 data.head()
```

Out[21]:

	age	job	marital	education_qual	call_type	day	mon	dur	num_calls	prev_outcome	y
0	58	8	married	tertiary	unknown	5	may	261	1	unknown	0
1	44	5	single	secondary	unknown	5	may	151	1	unknown	0
2	33	2	married	secondary	unknown	5	may	76	1	unknown	0
3	47	1	married	secondary	unknown	5	may	92	1	unknown	0
4	33	1	single	secondary	unknown	5	may	198	1	unknown	0

In [22]:

```
1 #Marital column encoding
2 marital_cat_ord = {'married':1, 'divorced':2, 'single':3}
3 data['marital'] = data['marital'].map(marital_cat_ord)
4 data.head()
```

Out[22]:

	age	job	marital	education_qual	call_type	day	mon	dur	num_calls	prev_outcome	y
0	58	8	1	tertiary	unknown	5	may	261	1	unknown	0
1	44	5	3	secondary	unknown	5	may	151	1	unknown	0
2	33	2	1	secondary	unknown	5	may	76	1	unknown	0
3	47	1	1	secondary	unknown	5	may	92	1	unknown	0
4	33	1	3	secondary	unknown	5	may	198	1	unknown	0

In [23]:

```
1 #education_qual column encoding
2 education_qual_cat_ord = {'primary':1, 'secondary':2, 'tertiary':3}
3 data['education_qual'] = data['education_qual'].map(education_qual_cat_ord)
4 data.head()
```

Out[23]:

	age	job	marital	education_qual	call_type	day	mon	dur	num_calls	prev_outcome	y
0	58	8	1	3	unknown	5	may	261	1	unknown	0
1	44	5	3	2	unknown	5	may	151	1	unknown	0
2	33	2	1	2	unknown	5	may	76	1	unknown	0
3	47	1	1	2	unknown	5	may	92	1	unknown	0
4	33	1	3	2	unknown	5	may	198	1	unknown	0

In [24]:

```
1 #call_type encoding
2 call_type_cat_ord = {'unknown':1, 'telephone':2, 'cellular':3}
3 data['call_type'] = data['call_type'].map(call_type_cat_ord)
4 data.head()
```

Out[24]:

	age	job	marital	education_qual	call_type	day	mon	dur	num_calls	prev_outcome	y
0	58	8	1	3	1	5	may	261	1	unknown	0
1	44	5	3	2	1	5	may	151	1	unknown	0
2	33	2	1	2	1	5	may	76	1	unknown	0
3	47	1	1	2	1	5	may	92	1	unknown	0
4	33	1	3	2	1	5	may	198	1	unknown	0

In [25]:

```

1 #month column encoding
2 mon_cat_ord = {'may':1, 'jul':2, 'jan':3, 'nov':4, 'jun':5, 'aug':6, 'feb':7, 'apr':8}
3 data['mon'] = data['mon'].map(mon_cat_ord)
4 data.head()

```

Out[25]:

	age	job	marital	education_qual	call_type	day	mon	dur	num_calls	prev_outcome	y
0	58	8	1	3	1	5	1	261	1	unknown	0
1	44	5	3	2	1	5	1	151	1	unknown	0
2	33	2	1	2	1	5	1	76	1	unknown	0
3	47	1	1	2	1	5	1	92	1	unknown	0
4	33	1	3	2	1	5	1	198	1	unknown	0

In [26]:

```

1 #prev_outcome encoding
2 prev_outcome_cat_ord = {'unknown':1, 'failure':2, 'other':3, 'success':4}
3 data['prev_outcome'] = data['prev_outcome'].map(prev_outcome_cat_ord)
4 data.head()

```

Out[26]:

	age	job	marital	education_qual	call_type	day	mon	dur	num_calls	prev_outcome	y
0	58	8	1	3	1	5	1	261	1	1	0
1	44	5	3	2	1	5	1	151	1	1	0
2	33	2	1	2	1	5	1	76	1	1	0
3	47	1	1	2	1	5	1	92	1	1	0
4	33	1	3	2	1	5	1	198	1	1	0

Data Splitting

In [27]:

```

1 #feature selection
2 X = data.drop('y', axis=1).values
3 y = data['y']
4

```

In [28]:

```
1 #data splitting
2 from sklearn.model_selection import train_test_split
3 X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size = 0.
```

In [29]:

```
1 #Scaling
2 from sklearn.preprocessing import StandardScaler
3 sc = StandardScaler()
4 X_train = sc.fit_transform(X_train)
5 X_test = sc.transform(X_test)
```

Linear Regression Model

In [30]:

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.metrics import confusion_matrix
3 from sklearn.metrics import roc_auc_score
4 from sklearn.metrics import RocCurveDisplay
5 from sklearn.metrics import classification_report
6 lr = LogisticRegression()
7 lr.fit(X_train, y_train)
```

Out[30]:

```
▼ LogisticRegression
LogisticRegression()
```

In [31]:

```
1 print("Confusion Matrix of LogisticRegression Model\n",confusion_matrix(y_test, lr.p
```

Confusion Matrix of LogisticRegression Model

```
[[9767  213]
 [ 942  380]]
```

In [32]:

```
1 print("ROC AUC Score of LogisticRegression Model ",roc_auc_score(y_test, lr.predict_
```


ROC AUC Score of LogisticRegression Model 0.8950860874547886

In [33]:

```
1 RocCurveDisplay.from_predictions(y_test, lr.predict_proba(X_test)[:, 1])
```

Out[33]:

<sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x1e190640790>



In [34]:

```
1 print("Classification of LogisticRegression Model \n\n",classification_report(lr.pre
```

Classification of LogisticRegression Model

	precision	recall	f1-score	support
0	0.98	0.91	0.94	10709
1	0.29	0.64	0.40	593
accuracy			0.90	11302
macro avg	0.63	0.78	0.67	11302
weighted avg	0.94	0.90	0.92	11302

K - Nearest Neighbors

In [35]:

```

1 from sklearn.neighbors import KNeighborsClassifier
2 from sklearn.model_selection import cross_val_score
3
4 knn = KNeighborsClassifier()
5 k_range = range(1, 21)
6 train_scores = []
7 k_scores = []
8
9 for k in k_range:
10     knn = KNeighborsClassifier(n_neighbors=k)
11     knn.fit(X_train, y_train)
12     train_score = roc_auc_score(y_train, knn.predict_proba(X_train)[: , 1])
13     val_score = np.mean(cross_val_score(knn, X_train, y_train, cv=5, scoring='roc_auc'))
14     k_score = val_score
15     print(f'K = {k}----Train Score = {train_score}----CV Score = {k_score}')

```

```

K = 1----Train Score = 1.0----CV Score = 0.6805474698080379
K = 2----Train Score = 0.9816713012632141----CV Score = 0.7545456587133581
K = 3----Train Score = 0.971532614675024----CV Score = 0.7900891782259667
K = 4----Train Score = 0.964362975271538----CV Score = 0.815826845106219
K = 5----Train Score = 0.9586269041883082----CV Score = 0.8325778494299222
K = 6----Train Score = 0.9544433047791259----CV Score = 0.8420307730347399
K = 7----Train Score = 0.9510942013662421----CV Score = 0.85091097292099
K = 8----Train Score = 0.9478977764108679----CV Score = 0.8572594231479057
K = 9----Train Score = 0.945291476871523----CV Score = 0.8620455527407822
K = 10----Train Score = 0.9430503416952553----CV Score = 0.866550361264429
5
K = 11----Train Score = 0.941185843975761----CV Score = 0.8698982224445725
K = 12----Train Score = 0.9396773279340771----CV Score = 0.872884588214910
2
K = 13----Train Score = 0.9379849434612679----CV Score = 0.874988773227901
9
K = 14----Train Score = 0.93666638816872----CV Score = 0.8782671667239603
K = 15----Train Score = 0.9355377388912832----CV Score = 0.880913435495020
4
K = 16----Train Score = 0.9345501855096098----CV Score = 0.882922813229577
4
K = 17----Train Score = 0.9335165797613852----CV Score = 0.884763881156378
2
K = 18----Train Score = 0.9325565407530351----CV Score = 0.885879488356666
9
K = 19----Train Score = 0.9319017534019638----CV Score = 0.886471757992637
1
K = 20----Train Score = 0.9310248385363105----CV Score = 0.887503965774054

```

In [36]:

```

1 knn = KNeighborsClassifier(n_neighbors=11)
2 knn.fit(X_train, y_train)

```

Out[36]:

```

▼      KNeighborsClassifier
KNeighborsClassifier(n_neighbors=11)

```

In [37]:

```
1 print("Confusion Matrix of KNeighborsClassifier\n",confusion_matrix(y_test, knn.pred
```

Confusion Matrix of KNeighborsClassifier

```
[[9708 272]
 [ 856 466]]
```

In [38]:

```
1 print("ROC AUC Score of KNeighborsClassifier",roc_auc_score(y_test, knn.predict_prob
```

ROC AUC Score of KNeighborsClassifier 0.8744125922040753

In [39]:

```
1 print("Classification of KNeighborsClassifier\n\n",classification_report(y_test, knn
```

Classification of KNeighborsClassifier

	precision	recall	f1-score	support
0	0.92	0.97	0.95	9980
1	0.63	0.35	0.45	1322
accuracy			0.90	11302
macro avg	0.78	0.66	0.70	11302
weighted avg	0.89	0.90	0.89	11302

DecisionTree

In [40]:

```

1 from sklearn.tree import DecisionTreeClassifier
2 dt = DecisionTreeClassifier()
3 for depth in [5,6,7,8,9,10,11,12,13,14,15,20,30,40]:
4     dt = DecisionTreeClassifier(max_depth=depth)
5     dt.fit(X_train, y_train)
6     train_score = roc_auc_score(y_train, dt.predict(X_train))
7     dt = DecisionTreeClassifier(max_depth=depth)
8     val_score = np.mean(cross_val_score(dt, X_train, y_train, cv=10, scoring='roc_auc'))
9     print("Depth : ", depth, " Training Accuracy : ", train_score, " Cross val score : ", val_score)

```

```

Depth : 5 Training Accuracy : 0.6784115498830073 Cross val score : 0.8605213774098273
Depth : 6 Training Accuracy : 0.713595204261992 Cross val score : 0.8770401773186407
Depth : 7 Training Accuracy : 0.7206457087446582 Cross val score : 0.879032460888698
Depth : 8 Training Accuracy : 0.7340240306957843 Cross val score : 0.8749398741296824
Depth : 9 Training Accuracy : 0.7624361893895616 Cross val score : 0.8617467247876182
Depth : 10 Training Accuracy : 0.7707150853844053 Cross val score : 0.8472624865897048
Depth : 11 Training Accuracy : 0.7925365306671541 Cross val score : 0.8207173280681864
Depth : 12 Training Accuracy : 0.8140482108407187 Cross val score : 0.7960270178218488
Depth : 13 Training Accuracy : 0.831992735666523 Cross val score : 0.7692698980034427
Depth : 14 Training Accuracy : 0.852242472370601 Cross val score : 0.7535395297627792
Depth : 15 Training Accuracy : 0.8816992972959573 Cross val score : 0.7471155627879238
Depth : 20 Training Accuracy : 0.9660589385075187 Cross val score : 0.7092734335594552
Depth : 30 Training Accuracy : 0.9996218805142425 Cross val score : 0.6975639961415118
Depth : 40 Training Accuracy : 1.0 Cross val score : 0.6999121537061563

```

In [41]:

```

1 dt = DecisionTreeClassifier(max_depth=11)
2 dt.fit(X_train, y_train)

```

Out[41]:

```

DecisionTreeClassifier
DecisionTreeClassifier(max_depth=11)

```

In [42]:

```

1 print("Confusion Matrix of DecisionTreeClassifier\n", confusion_matrix(y_test, dt.predict(X_test)))

```

```

Confusion Matrix of DecisionTreeClassifier
[[9519  461]
 [ 690  632]]

```


In [43]:

```
1 print("ROC AUC Score of DecisionTreeClassifier",roc_auc_score(y_test, dt.predict_proba(X_test)))
```

ROC AUC Score of DecisionTreeClassifier 0.837099046807685

In [44]:

```
1 print("Classification of DecisionTreeClassifier\n\n",classification_report(y_test, dt.predict_proba(X_test)))
```

Classification of DecisionTreeClassifier

	precision	recall	f1-score	support
0	0.93	0.95	0.94	9980
1	0.58	0.48	0.52	1322
accuracy			0.90	11302
macro avg	0.76	0.72	0.73	11302
weighted avg	0.89	0.90	0.89	11302

In []:

```
1
```