

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```


In [2]:

```
1 # You have to include the full link to the csv file containing your dataset
2 employee_df = pd.read_csv('D:/Data Science for Business Package/1. Human Resources D
3 employee_df
```

Out[2]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education
0	41	Yes	Travel_Rarely	1102	Sales	1	2
1	49	No	Travel_Frequently	279	Research & Development	8	1
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2
3	33	No	Travel_Frequently	1392	Research & Development	3	4
4	27	No	Travel_Rarely	591	Research & Development	2	1
...	...	...	...	...	...	...	..
1465	36	No	Travel_Frequently	884	Research & Development	23	2
1466	39	No	Travel_Rarely	613	Research & Development	6	1
1467	27	No	Travel_Rarely	155	Research & Development	4	3
1468	49	No	Travel_Frequently	1023	Sales	2	3
1469	34	No	Travel_Rarely	628	Research & Development	8	3

1470 rows × 35 columns



In [3]:

```
1 print(employee_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                    1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EmployeeNumber                      1470 non-null   int64
10  EnvironmentSatisfaction              1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                          1470 non-null   int64
13  JobInvolvement                      1470 non-null   int64
14  JobLevel                            1470 non-null   int64
15  JobRole                             1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                      1470 non-null   object
18  MonthlyIncome                      1470 non-null   int64
19  MonthlyRate                         1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                             1470 non-null   object
22  OverTime                           1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                      1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
None
```

In [4]:

```

1 print("Attrition columns record values",employee_df['Attrition'].unique())
2 print("Overtime columns record values",employee_df['OverTime'].unique())
3 print("Over18 columns record values",employee_df['Over18'].unique())
4
5 #categorical data into numerical transformation manually
6 employee_df['Attrition'] = employee_df['Attrition'].apply(lambda x: 1 if x == 'Yes'
7 employee_df['OverTime'] = employee_df['OverTime'].apply(lambda x: 1 if x == 'Yes' el
8 employee_df['Over18'] = employee_df['Over18'].apply(lambda x: 1 if x == 'Y' else 0)
9

```

Attrition columns record values ['Yes' 'No']  
 Overtime columns record values ['Yes' 'No']  
 Over18 columns record values ['Y']

In [5]:

```

1 #checking for null values
2 employee_df.isnull().sum()
3

```

Out[5]:

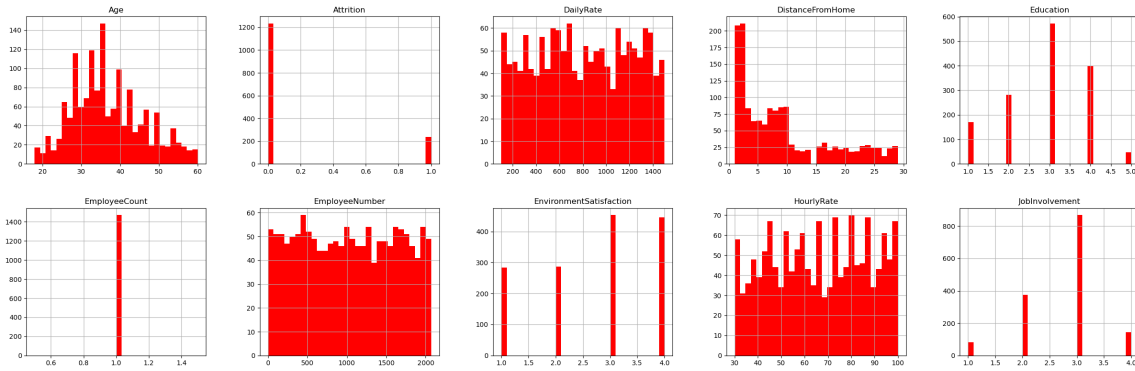
Age	0
Attrition	0
BusinessTravel	0
DailyRate	0
Department	0
DistanceFromHome	0
Education	0
EducationField	0
EmployeeCount	0
EmployeeNumber	0
EnvironmentSatisfaction	0
Gender	0
HourlyRate	0
JobInvolvement	0
JobLevel	0
JobRole	0
JobSatisfaction	0
MaritalStatus	0
MonthlyIncome	0
MonthlyRate	0
NumCompaniesWorked	0
Over18	0
OverTime	0
PercentSalaryHike	0
PerformanceRating	0
RelationshipSatisfaction	0
StandardHours	0
StockOptionLevel	0
TotalWorkingYears	0
TrainingTimesLastYear	0
WorkLifeBalance	0
YearsAtCompany	0
YearsInCurrentRole	0
YearsSinceLastPromotion	0
YearsWithCurrManager	0
dtype: int64	

In [6]:

```
1 employee_df.hist(bins = 30,figsize = (30,30), color = 'r')
```

Out[6]:

```
array([[<AxesSubplot: title={'center': 'Age'}>,
        <AxesSubplot: title={'center': 'Attrition'}>,
        <AxesSubplot: title={'center': 'DailyRate'}>,
        <AxesSubplot: title={'center': 'DistanceFromHome'}>,
        <AxesSubplot: title={'center': 'Education'}>],
       [<AxesSubplot: title={'center': 'EmployeeCount'}>,
        <AxesSubplot: title={'center': 'EmployeeNumber'}>,
        <AxesSubplot: title={'center': 'EnvironmentSatisfaction'}>,
        <AxesSubplot: title={'center': 'HourlyRate'}>,
        <AxesSubplot: title={'center': 'JobInvolvement'}>],
       [<AxesSubplot: title={'center': 'JobLevel'}>,
        <AxesSubplot: title={'center': 'JobSatisfaction'}>,
        <AxesSubplot: title={'center': 'MonthlyIncome'}>,
        <AxesSubplot: title={'center': 'MonthlyRate'}>,
        <AxesSubplot: title={'center': 'NumCompaniesWorked'}>],
       [<AxesSubplot: title={'center': 'Over18'}>,
        <AxesSubplot: title={'center': 'OverTime'}>,
        <AxesSubplot: title={'center': 'PercentSalaryHike'}>,
        <AxesSubplot: title={'center': 'PerformanceRating'}>,
        <AxesSubplot: title={'center': 'RelationshipSatisfaction'}>],
       [<AxesSubplot: title={'center': 'StandardHours'}>,
        <AxesSubplot: title={'center': 'StockOptionLevel'}>,
        <AxesSubplot: title={'center': 'TotalWorkingYears'}>,
        <AxesSubplot: title={'center': 'TrainingTimesLastYear'}>,
        <AxesSubplot: title={'center': 'WorkLifeBalance'}>],
       [<AxesSubplot: title={'center': 'YearsAtCompany'}>,
        <AxesSubplot: title={'center': 'YearsInCurrentRole'}>,
        <AxesSubplot: title={'center': 'YearsSinceLastPromotion'}>,
        <AxesSubplot: title={'center': 'YearsWithCurrManager'}>,
        <AxesSubplot: >]], dtype=object)
```



```

1 # From the above visualization we can drop 'EmployeeCount', 'StandardHours' and 'OverTime'
2 # Let's drop 'EmployeeNumber' as well it is 10
3 employee_df.drop(['EmployeeCount', 'StandardHours', 'OverTime', 'EmployeeNumber'], axis=1)


```

```

1 print("People who left company {} and their percentage {:.2f}".format(len(employee_df[Attrition==1]),
2 print("People who stayed in company {} and their percentage {:.2f}".format(len(employee_df[Attrition==0]),

```

People who left company 237 and their percentage 16.12  
 People who stayed in company 1233 and their percentage 83.88



In [9]:

```

1 #correlation plot
2
3 correlations = employee_df.corr()
4 fig, ax = plt.subplots(figsize = (20, 20))
5 sns.heatmap(correlations, annot = True)
6


```

C:\Users\Pradeep\AppData\Local\Temp\ipykernel\_22596\3969665670.py:3: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
correlations = employee_df.corr()
```

Out[9]:

&lt;AxesSubplot: &gt;



In [10]:


```

1 #visualizing age of people who stayed and left the company
2 plt.figure(figsize=[20,10])
3 sns.countplot(x = 'Age', hue = 'Attrition', data = employee_df)

```

Out[10]:

&lt;AxesSubplot: xlabel='Age', ylabel='count'&gt;




In [11]:

```
1 employee_df.describe()
```

Out[11]:

	Age	Attrition	DailyRate	DistanceFromHome	Education	Environment
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.000000	
mean	36.923810	0.161224	802.485714	9.192517	2.912925	
std	9.135373	0.367863	403.509100	8.106864	1.024165	
min	18.000000	0.000000	102.000000	1.000000	1.000000	
25%	30.000000	0.000000	465.000000	2.000000	2.000000	
50%	36.000000	0.000000	802.000000	7.000000	3.000000	
75%	43.000000	0.000000	1157.000000	14.000000	4.000000	
max	60.000000	1.000000	1499.000000	29.000000	5.000000	

8 rows × 25 columns



In [12]:

```
1 #encoding for categorical column  
2 employee_df = pd.get_dummies(employee_df, columns=['Department', 'BusinessTravel', 'Education'])  
3 employee_df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1470 entries, 0 to 1469
```

```
Data columns (total 51 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	1470 non-null	int64
1	Attrition	1470 non-null	int64
2	DailyRate	1470 non-null	int64
3	DistanceFromHome	1470 non-null	int64
4	Education	1470 non-null	int64
5	EnvironmentSatisfaction	1470 non-null	int64
6	HourlyRate	1470 non-null	int64
7	JobInvolvement	1470 non-null	int64
8	JobLevel	1470 non-null	int64
9	JobSatisfaction	1470 non-null	int64
10	MonthlyIncome	1470 non-null	int64
11	MonthlyRate	1470 non-null	int64
12	NumCompaniesWorked	1470 non-null	int64
13	OverTime	1470 non-null	int64
14	PercentSalaryHike	1470 non-null	int64
15	PerformanceRating	1470 non-null	int64
16	RelationshipSatisfaction	1470 non-null	int64
17	StockOptionLevel	1470 non-null	int64
18	TotalWorkingYears	1470 non-null	int64
19	TrainingTimesLastYear	1470 non-null	int64
20	WorkLifeBalance	1470 non-null	int64
21	YearsAtCompany	1470 non-null	int64
22	YearsInCurrentRole	1470 non-null	int64
23	YearsSinceLastPromotion	1470 non-null	int64
24	YearsWithCurrManager	1470 non-null	int64
25	Department_Human Resources	1470 non-null	uint8
26	Department_Research & Development	1470 non-null	uint8
27	Department_Sales	1470 non-null	uint8
28	BusinessTravel_Non-Travel	1470 non-null	uint8
29	BusinessTravel_Travel_Frequently	1470 non-null	uint8
30	BusinessTravel_Travel_Rarely	1470 non-null	uint8
31	EducationField_Human Resources	1470 non-null	uint8
32	EducationField_Life Sciences	1470 non-null	uint8
33	EducationField_Marketing	1470 non-null	uint8
34	EducationField_Medical	1470 non-null	uint8
35	EducationField_Other	1470 non-null	uint8
36	EducationField_Technical Degree	1470 non-null	uint8
37	Gender_Female	1470 non-null	uint8
38	Gender_Male	1470 non-null	uint8
39	JobRole_Healthcare Representative	1470 non-null	uint8
40	JobRole_Human Resources	1470 non-null	uint8
41	JobRole_Laboratory Technician	1470 non-null	uint8
42	JobRole_Manager	1470 non-null	uint8
43	JobRole_Manufacturing Director	1470 non-null	uint8
44	JobRole_Research Director	1470 non-null	uint8
45	JobRole_Research Scientist	1470 non-null	uint8
46	JobRole_Sales Executive	1470 non-null	uint8
47	JobRole_Sales Representative	1470 non-null	uint8
48	MaritalStatus_Divorced	1470 non-null	uint8
49	MaritalStatus_Married	1470 non-null	uint8
50	MaritalStatus_Single	1470 non-null	uint8

```
dtypes: int64(25), uint8(26)
```

```
memory usage: 324.6 KB
```

In [13]:

```
1 # X and y variable
2 X = employee_df.drop(['Attrition'],axis=1)
3 y = employee_df['Attrition']
```

In [14]:

```
1 #scaling
2 from sklearn.preprocessing import MinMaxScaler
3 scaler = MinMaxScaler()
4 X = scaler.fit_transform(X)
```

In [15]:

```
1 #splitting data
2 from sklearn.model_selection import train_test_split
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25)
4 X_train.shape
```

Out[15]:

(1102, 50)

In [16]:

```
1 # LogisticRegression ML training
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.metrics import accuracy_score
4
5 model = LogisticRegression()
6 model.fit(X_train, y_train)
7
8 y_pred = model.predict(X_test)
```

In [17]:

```
1 # comparing predicted vs actual value
2 data_pred = pd.DataFrame(y_pred,columns=['Predicted Value'])
3 data_pred['Actual Value'] = y_test.values
4 data_pred
```

Out[17]:

	Predicted Value	Actual Value
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...	...	...
363	0	1
364	0	0
365	0	0
366	0	1
367	0	0

368 rows × 2 columns

In [18]:


```

1 #Model Evaluation
2 from sklearn.metrics import confusion_matrix, classification_report
3
4 print("Accuracy {:.2f} %".format(accuracy_score(y_pred, y_test)*100))
5 # Testing Set Performance
6 cm = confusion_matrix(y_pred, y_test)
7 sns.heatmap(cm, annot=True)
8
9 print('\n',classification_report(y_test, y_pred))
10

```

Accuracy 86.68 %

	precision	recall	f1-score	support
0	0.89	0.96	0.92	306
1	0.69	0.39	0.49	62
accuracy			0.87	368
macro avg	0.79	0.68	0.71	368
weighted avg	0.85	0.87	0.85	368





In [21]:

```
1 y_pred = model.predict(X_test)
2 y_pred = y_pred>0.5
3 y_pred
4
```

```
12/12 [=====] - 0s 4ms/step
```

Out[21]:


[illegible]

In [22]:

```
1 plt.plot(epochs_hist.history['loss'])
2 plt.title('Model Loss')
3 plt.xlabel('Epoch')
4 plt.ylabel('Training Loss')
5 plt.legend(['Training Loss'])
```

Out[22]:

<matplotlib.legend.Legend at 0x1c9e6267310>




In [23]:

```
1 plt.plot(epochs_hist.history['accuracy'])
2 plt.title('Model Accuracy')
3 plt.xlabel('Epoch')
4 plt.ylabel('Training Accuracy')
5 plt.legend(['Training Accuracy'])
```

Out[23]:

&lt;matplotlib.legend.Legend at 0x1c9e6468d90&gt;






In [24]:

```
1 #confusion matrix
2 cm = confusion_matrix(y_test, y_pred)
3 sns.heatmap(cm, annot=True)
4
```

Out[24]:

&lt;AxesSubplot: &gt;



In [25]:

```
1 print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.94	0.91	306
1	0.57	0.39	0.46	62
accuracy			0.85	368
macro avg	0.73	0.66	0.69	368
weighted avg	0.83	0.85	0.84	368

In [ ]:

```
1
```