

In [1]:

```

1 #importing data
2
3 import pandas as pd
4 insurance = pd.read_csv("D:/Data Science and Deep Learning for Business/datasciencef
5 insurance.head()
```

Out[1]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [2]:

```

1 #describing dataset
2
3 print ("Rows      : " , insurance.shape[0])
4 print ("Columns   : " , insurance.shape[1])
5 print ("\nFeatures : \n" , insurance.columns.tolist())
6 print ("\nMissing values : \n", insurance.isnull().sum())
7 print ("\nUnique values : \n",insurance.nunique())
```

```

Rows      : 1338
Columns   : 7
```

```

Features :
['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges']
```

```

Missing values :
age          0
sex          0
bmi          0
children     0
smoker       0
region       0
charges      0
dtype: int64
```

```

Unique values :
age          47
sex          2
bmi         548
children     6
smoker       2
region       4
charges     1337
dtype: int64
```

In [3]:

```
1 #correlation
2 insurance.corr()
```

C:\Users\Pradeep\AppData\Local\Temp\ipykernel\_23608\641068510.py:2: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
insurance.corr()
```

Out[3]:

	age	bmi	children	charges
age	1.000000	0.109272	0.042469	0.299008
bmi	0.109272	1.000000	0.012759	0.198341
children	0.042469	0.012759	1.000000	0.067998
charges	0.299008	0.198341	0.067998	1.000000

In [4]:

```


1 #correlation plot
2 import matplotlib.pyplot as plt
3
4 def plot_corr(df):
5     corr = df.corr()
6     fig, ax = plt.subplots(figsize=(10, 10))
7     ax.legend()
8     cax = ax.matshow(corr) #array as matrix
9     fig.colorbar(cax)
10    plt.xticks(range(len(corr.columns)), corr.columns, rotation='vertical')
11    plt.yticks(range(len(corr.columns)), corr.columns)
12
13 plot_corr(insurance)

```

C:\Users\Pradeep\AppData\Local\Temp\ipykernel\_23608\4274883086.py:5: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
corr = df.corr()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.







In [5]:

```
1 insurance.hist(bins=100,figsize=(20,20))
```

Out[5]:

```
array([[<AxesSubplot: title={'center': 'age'}>,  
       <AxesSubplot: title={'center': 'bmi'}>],  
       [<AxesSubplot: title={'center': 'children'}>,  
       <AxesSubplot: title={'center': 'charges'}>]], dtype=object)
```




In [6]:

```
1 insurance['sex'].value_counts().plot(kind='bar')
```

Out[6]:

<AxesSubplot: >




In [7]:

```
1 insurance['smoker'].value_counts().plot(kind='bar')
```


Out[7]:

&lt;AxesSubplot: &gt;



In [8]:

```
1 fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(15, 5))
2 insurance.plot(kind='scatter', x='age', y='charges', alpha=0.5, color='green', ax=axes[0])
3 insurance.plot(kind='scatter', x='bmi', y='charges', alpha=0.5, color='red', ax=axes[1])
4 insurance.plot(kind='scatter', x='children', y='charges', alpha=0.5, color='blue', ax=axes[2])
5 plt.show()
```




In [9]:

```
1 # Importing Seaborn library
2 import seaborn as sns
3 sns.scatterplot(x="bmi", y="charges", data=insurance, hue='smoker')
```

Out[9]:

&lt;AxesSubplot: xlabel='bmi', ylabel='charges'&gt;



In [10]:

```
1 insurance.head()
```

Out[10]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [11]:

```
1 insurance['region'].unique()
```

Out[11]:

array(['southwest', 'southeast', 'northwest', 'northeast'], dtype=object)

In [12]:

```
1 insurance.drop(["region"], axis=1, inplace=True)
2 insurance.head()
```

Out[12]:

	age	sex	bmi	children	smoker	charges
0	19	female	27.900	0	yes	16884.92400
1	18	male	33.770	1	no	1725.55230
2	28	male	33.000	3	no	4449.46200
3	33	male	22.705	0	no	21984.47061
4	32	male	28.880	0	no	3866.85520

In [13]:

```
1 insurance['sex'].unique()
```

Out[13]:

```
array(['female', 'male'], dtype=object)
```

In [14]:

```
1 insurance['sex'] = insurance['sex'].map(lambda s :1 if s == 'female' else 0)
2 insurance.head()
```

Out[14]:

	age	sex	bmi	children	smoker	charges
0	19	1	27.900	0	yes	16884.92400
1	18	0	33.770	1	no	1725.55230
2	28	0	33.000	3	no	4449.46200
3	33	0	22.705	0	no	21984.47061
4	32	0	28.880	0	no	3866.85520

In [15]:

```
1 insurance['smoker'].unique()
```

Out[15]:

```
array(['yes', 'no'], dtype=object)
```



In [16]:

```
1 insurance['smoker'] = insurance['smoker'].map(lambda s :1 if s == 'yes' else 0)
2 insurance.head()
```

Out[16]:

	age	sex	bmi	children	smoker	charges
0	19	1	27.900	0	1	16884.92400
1	18	0	33.770	1	0	1725.55230
2	28	0	33.000	3	0	4449.46200
3	33	0	22.705	0	0	21984.47061
4	32	0	28.880	0	0	3866.85520

In [17]:

```
1 X = insurance.drop(['charges'], axis = 1)
2 y = insurance['charges']
```

In [18]:

```
1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LinearRegression
3
4 X_train, X_test, y_train, y_test = train_test_split(X, y)
5 lr = LinearRegression().fit(X_train, y_train)
6
7 y_train_pred = lr.predict(X_train)
8 y_test_pred = lr.predict(X_test)
9
10 print(lr.score(X_test, y_test))
```

0.7608249670500205

In [19]:

```
1 results = pd.DataFrame({'Actual': y_test, 'Predicted': y_test_pred})
2 results
```

Out[19]:

	Actual	Predicted
732	4234.92700	5553.750604
708	6113.23105	7507.860180
766	8062.76400	10758.101412
1281	24535.69855	33420.147504
1000	17361.76610	27344.626895
...	...	...
911	33732.68670	25531.328016
1051	14394.55790	12859.213123
1041	1704.70015	-197.426943
482	1622.18850	2481.649508
1164	7153.55390	8243.460729

335 rows × 2 columns

In [20]:

```
1 # Normalize the data
2 from sklearn.preprocessing import StandardScaler
3
4 sc = StandardScaler()
5 X_train = sc.fit_transform(X_train)
6 X_test = sc.transform(X_test)
7
8 pd.DataFrame(X_train).head()
```

Out[20]:

	0	1	2	3	4
0	-1.064757	-0.991067	-0.109293	-0.921510	2.029125
1	1.697735	-0.991067	1.798254	-0.921510	-0.492823
2	-1.064757	-0.991067	-1.121243	-0.921510	-0.492823
3	0.281073	1.009014	-0.097821	-0.093307	-0.492823
4	0.847738	1.009014	1.185348	-0.093307	-0.492823

In [21]:

```
1 pd.DataFrame(y_train).head()
```

Out[21]:

	charges
1250	18648.42170
170	13405.39030
693	2352.96845
1263	7337.74800
147	9877.60770

In [22]:

```
1 #Linear Regression model
2 from sklearn.linear_model import LinearRegression
3
4 multiple_linear_reg = LinearRegression(fit_intercept=False)
5 multiple_linear_reg.fit(X_train, y_train)
```

Out[22]:

	LinearRegression
LinearRegression(fit_intercept=False)	

In [23]:

```
1 #PolynomialFeatures
2 from sklearn.preprocessing import PolynomialFeatures
3
4 polynomial_features = PolynomialFeatures(degree=3)
5 x_train_poly = polynomial_features.fit_transform(X_train)
6 x_test_poly = polynomial_features.fit_transform(X_test)
7
8 polynomial_reg = LinearRegression(fit_intercept=False)
9 polynomial_reg.fit(x_train_poly, y_train)
```

Out[23]:

	LinearRegression
LinearRegression(fit_intercept=False)	

In [24]:

```
1 #Decision Tree Regression model
2 from sklearn.tree import DecisionTreeRegressor
3
4 decision_tree_reg = DecisionTreeRegressor(max_depth=5, random_state=13)
5 decision_tree_reg.fit(X_train, y_train)
```

Out[24]:

```
DecisionTreeRegressor
DecisionTreeRegressor(max_depth=5, random_state=13)
```

In [25]:

```
1 #Random Forest Regression model
2 from sklearn.ensemble import RandomForestRegressor
3
4 random_forest_reg = RandomForestRegressor(n_estimators=400, max_depth=5, random_state=13)
5 random_forest_reg.fit(X_train, y_train)
```

Out[25]:

```
RandomForestRegressor
RandomForestRegressor(max_depth=5, n_estimators=400, random_state=13)
```

In [26]:

```
1 #SVR model
2
3 from sklearn.svm import SVR
4
5 support_vector_reg = SVR(gamma="auto", kernel="linear", C=1000)
6 support_vector_reg.fit(X_train, y_train)
```

Out[26]:

```
SVR
SVR(C=1000, gamma='auto', kernel='linear')
```

In [27]:

```
1 #Importing evaluation metrics
2
3 from sklearn.model_selection import cross_val_predict
4 from sklearn.metrics import r2_score
5 from sklearn.metrics import mean_squared_error
6 from math import sqrt
```

In [28]:

```

1  #Evaluating Multiple Linear Regression Model
2
3  # Prediction with training dataset:
4  y_pred_MLR_train = multiple_linear_reg.predict(X_train)
5
6  # Prediction with testing dataset:
7  y_pred_MLR_test = multiple_linear_reg.predict(X_test)
8
9  # Find training accuracy for this model:
10 accuracy_MLR_train = r2_score(y_train, y_pred_MLR_train)
11 print("Training Accuracy for Multiple Linear Regression Model: ", accuracy_MLR_train)
12
13 # Find testing accuracy for this model:
14 accuracy_MLR_test = r2_score(y_test, y_pred_MLR_test)
15 print("Testing Accuracy for Multiple Linear Regression Model: ", accuracy_MLR_test)
16
17 # Find RMSE for training data:
18 RMSE_MLR_train = sqrt(mean_squared_error(y_train, y_pred_MLR_train))
19 print("RMSE for Training Data: ", RMSE_MLR_train)
20
21 # Find RMSE for testing data:
22 RMSE_MLR_test = sqrt(mean_squared_error(y_test, y_pred_MLR_test))
23 print("RMSE for Testing Data: ", RMSE_MLR_test)
24
25 # Prediction with 10-Fold Cross Validation:
26 y_pred_cv_MLR = cross_val_predict(multiple_linear_reg, X, y, cv=10)
27
28 # Find accuracy after 10-Fold Cross Validation
29 accuracy_cv_MLR = r2_score(y, y_pred_cv_MLR)
30 print("Accuracy for 10-Fold Cross Predicted Multiple Linear Regression Model: ", acc

```

Training Accuracy for Multiple Linear Regression Model: -0.44214251208022

58

Testing Accuracy for Multiple Linear Regression Model: -0.308218253498038

2

RMSE for Training Data: 14178.647894150441

RMSE for Testing Data: 14728.890264369307

Accuracy for 10-Fold Cross Predicted Multiple Linear Regression Model: 0.

717113419200113

In [29]:

```
1 #Evaluating Polynomial Regression Model
2
3 # Prediction with training dataset:
4 y_pred_PR_train = polynomial_reg.predict(x_train_poly)
5
6 # Prediction with testing dataset:
7 y_pred_PR_test = polynomial_reg.predict(x_test_poly)
8
9 # Find training accuracy for this model:
10 accuracy_PR_train = r2_score(y_train, y_pred_PR_train)
11 print("Training Accuracy for Polynomial Regression Model: ", accuracy_PR_train)
12
13 # Find testing accuracy for this model:
14 accuracy_PR_test = r2_score(y_test, y_pred_PR_test)
15 print("Testing Accuracy for Polynomial Regression Model: ", accuracy_PR_test)
16
17 # Find RMSE for training data:
18 RMSE_PR_train = sqrt(mean_squared_error(y_train, y_pred_PR_train))
19 print("RMSE for Training Data: ", RMSE_PR_train)
20
21 # Find RMSE for testing data:
22 RMSE_PR_test = sqrt(mean_squared_error(y_test, y_pred_PR_test))
23 print("RMSE for Testing Data: ", RMSE_PR_test)
24
25 # Prediction with 10-Fold Cross Validation:
26 y_pred_cv_PR = cross_val_predict(polynomial_reg, polynomial_features.fit_transform(X_train), y_train, cv=10)
27
28 # Find accuracy after 10-Fold Cross Validation
29 accuracy_cv_PR = r2_score(y, y_pred_cv_PR)
30 print("Accuracy for 10-Fold Cross Predicted Polynomial Regression Model: ", accuracy_cv_PR)
```

Training Accuracy for Polynomial Regression Model: 0.8482275078787543

Testing Accuracy for Polynomial Regression Model: 0.8424785147329269

RMSE for Training Data: 4599.676205617919

RMSE for Testing Data: 5110.928711122368

Accuracy for 10-Fold Cross Predicted Polynomial Regression Model: 0.83910

72917688998

In [30]:

```
1 # Evaluating Decision Tree Regression Model
2
3 # Prediction with training dataset:
4 y_pred_DTR_train = decision_tree_reg.predict(X_train)
5
6 # Prediction with testing dataset:
7 y_pred_DTR_test = decision_tree_reg.predict(X_test)
8
9 # Find training accuracy for this model:
10 accuracy_DTR_train = r2_score(y_train, y_pred_DTR_train)
11 print("Training Accuracy for Decision Tree Regression Model: ", accuracy_DTR_train)
12
13 # Find testing accuracy for this model:
14 accuracy_DTR_test = r2_score(y_test, y_pred_DTR_test)
15 print("Testing Accuracy for Decision Tree Regression Model: ", accuracy_DTR_test)
16
17 # Find RMSE for training data:
18 RMSE_DTR_train = sqrt(mean_squared_error(y_train, y_pred_DTR_train))
19 print("RMSE for Training Data: ", RMSE_DTR_train)
20
21 # Find RMSE for testing data:
22 RMSE_DTR_test = sqrt(mean_squared_error(y_test, y_pred_DTR_test))
23 print("RMSE for Testing Data: ", RMSE_DTR_test)
24
25 # Prediction with 10-Fold Cross Validation:
26 y_pred_cv_DTR = cross_val_predict(decision_tree_reg, X, y, cv=10)
27
28 # Find accuracy after 10-Fold Cross Validation
29 accuracy_cv_DTR = r2_score(y, y_pred_cv_DTR)
30 print("Accuracy for 10-Fold Cross Predicted Decision Tree Regression Model: ", accur
```

Training Accuracy for Decision Tree Regression Model: 0.8810575984972877  
Testing Accuracy for Decision Tree Regression Model: 0.8312217664502496  
RMSE for Training Data: 4071.9185223433237  
RMSE for Testing Data: 5290.395531447086  
Accuracy for 10-Fold Cross Predicted Decision Tree Regression Model: 0.8494241031595924

In [31]:

```
1 # Evaluating Random Forest Regression Model
2
3 # Prediction with training dataset:
4 y_pred_RFR_train = random_forest_reg.predict(X_train)
5
6 # Prediction with testing dataset:
7 y_pred_RFR_test = random_forest_reg.predict(X_test)
8
9 # Find training accuracy for this model:
10 accuracy_RFR_train = r2_score(y_train, y_pred_RFR_train)
11 print("Training Accuracy for Random Forest Regression Model: ", accuracy_RFR_train)
12
13 # Find testing accuracy for this model:
14 accuracy_RFR_test = r2_score(y_test, y_pred_RFR_test)
15 print("Testing Accuracy for Random Forest Regression Model: ", accuracy_RFR_test)
16
17 # Find RMSE for training data:
18 RMSE_RFR_train = sqrt(mean_squared_error(y_train, y_pred_RFR_train))
19 print("RMSE for Training Data: ", RMSE_RFR_train)
20
21 # Find RMSE for testing data:
22 RMSE_RFR_test = sqrt(mean_squared_error(y_test, y_pred_RFR_test))
23 print("RMSE for Testing Data: ", RMSE_RFR_test)
24
25 # Prediction with 10-Fold Cross Validation:
26 y_pred_cv_RFR = cross_val_predict(random_forest_reg, X, y, cv=10)
27
28 # Find accuracy after 10-Fold Cross Validation
29 accuracy_cv_RFR = r2_score(y, y_pred_cv_RFR)
30 print("Accuracy for 10-Fold Cross Predicted Random Forest Regression Model: ", accur
```

Training Accuracy for Random Forest Regression Model: 0.8881034695908978  
Testing Accuracy for Random Forest Regression Model: 0.8709486381615694  
RMSE for Training Data: 3949.4719903815003  
RMSE for Testing Data: 4626.059562572722  
Accuracy for 10-Fold Cross Predicted Random Forest Regression Model: 0.8573788696785247



In [32]:

```
1 # Evaluating Support Vector Regression Model
2
3 # Prediction with training dataset:
4 y_pred_SVR_train = support_vector_reg.predict(X_train)
5
6 # Prediction with testing dataset:
7 y_pred_SVR_test = support_vector_reg.predict(X_test)
8
9 # Find training accuracy for this model:
10 accuracy_SVR_train = r2_score(y_train, y_pred_SVR_train)
11 print("Training Accuracy for Support Vector Regression Model: ", accuracy_SVR_train)
12
13 # Find testing accuracy for this model:
14 accuracy_SVR_test = r2_score(y_test, y_pred_SVR_test)
15 print("Testing Accuracy for Support Vector Regression Model: ", accuracy_SVR_test)
16
17 # Find RMSE for training data:
18 RMSE_SVR_train = sqrt(mean_squared_error(y_train, y_pred_SVR_train))
19 print("RMSE for Training Data: ", RMSE_SVR_train)
20
21 # Find RMSE for testing data:
22 RMSE_SVR_test = sqrt(mean_squared_error(y_test, y_pred_SVR_test))
23 print("RMSE for Testing Data: ", RMSE_SVR_test)
24
25 # Prediction with 10-Fold Cross Validation:
26 y_pred_cv_SVR = cross_val_predict(support_vector_reg, X, y, cv=10)
27
28 # Find accuracy after 10-Fold Cross Validation
29 accuracy_cv_SVR = r2_score(y, y_pred_cv_SVR)
30 print("Accuracy for 10-Fold Cross Predicted Support Vector Regression Model: ", accu
```

Training Accuracy for Support Vector Regression Model: 0.7114656153556187  
Testing Accuracy for Support Vector Regression Model: 0.7314092980873221  
RMSE for Training Data: 6342.04786792858  
RMSE for Testing Data: 6673.83417232676  
Accuracy for 10-Fold Cross Predicted Support Vector Regression Model: 0.7058131221977515

In [33]:

```

1 # Compare all results in one table
2 training_accuracies = [accuracy_MLR_train, accuracy_PR_train, accuracy_DTR_train, accuracy_RFR_train, accuracy_SVR_train]
3 testing_accuracies = [accuracy_MLR_test, accuracy_PR_test, accuracy_DTR_test, accuracy_RFR_test, accuracy_SVR_test]
4 training_RMSE = [RMSE_MLR_train, RMSE_PR_train, RMSE_DTR_train, RMSE_RFR_train, RMSE_SVR_train]
5 testing_RMSE = [RMSE_MLR_test, RMSE_PR_test, RMSE_DTR_test, RMSE_RFR_test, RMSE_SVR_test]
6 cv_accuracies = [accuracy_cv_MLR, accuracy_cv_PR, accuracy_cv_DTR, accuracy_cv_RFR, accuracy_cv_SVR]
7
8 parameters = ["fit_intercept=False", "fit_intercept=False", "max_depth=5", "n_estimators=400", "kernel='linear'", "C=1000"]
9
10 table_data = {"Parameters": parameters, "Training Accuracy": training_accuracies, "Testing Accuracy": testing_accuracies, "Training RMSE": training_RMSE, "Testing RMSE": testing_RMSE, "10-Fold Score": cv_accuracies}
11
12 model_names = ["Multiple Linear Regression", "Polynomial Regression", "Decision Tree Regression", "Random Forest Regression", "Support Vector Regression"]
13
14 table_dataframe = pd.DataFrame(data=table_data, index=model_names)
15 table_dataframe

```

Out[33]:

	Parameters	Training Accuracy	Testing Accuracy	Training RMSE	Testing RMSE	10-Fold Score
<b>Multiple Linear Regression</b>	fit_intercept=False	-0.442143	-0.308218	14178.647894	14728.890264	0.717113
<b>Polynomial Regression</b>	fit_intercept=False	0.848228	0.842479	4599.676206	5110.928711	0.839107
<b>Decision Tree Regression</b>	max_depth=5	0.881058	0.831222	4071.918522	5290.395531	0.849424
<b>Random Forest Regression</b>	n_estimators=400, max_depth=5	0.888103	0.870949	3949.471990	4626.059563	0.857379
<b>Support Vector Regression</b>	kernel="linear", C=1000	0.711466	0.731409	6342.047868	6673.834172	0.705813

In [34]:

```

1 #test on new input data
2 input_data = {'age': [35],
3               'sex': ['male'],
4               'bmi': [26],
5               'children': [0],
6               'smoker': ['no'],
7               'region': ['southeast']}
8
9 input_data = pd.DataFrame(input_data)
10 input_data

```

Out[34]:

	age	sex	bmi	children	smoker	region
0	35	male	26	0	no	southeast

In [35]:

```
1 #Input data pre-processing
2 input_data.drop(["region"], axis=1, inplace=True)
3 input_data['sex'] = input_data['sex'].map(lambda s :1 if s == 'female' else 0)
4 input_data['smoker'] = input_data['smoker'].map(lambda s :1 if s == 'yes' else 0)
5 input_data
```

Out[35]:

	age	sex	bmi	children	smoker
0	35	0	26	0	0

In [36]:

```
1 # Scale our input data
2 input_data = sc.transform(input_data)
3 input_data
```

Out[36]:

```
array([[ -0.28559244, -0.99106682, -0.73694802, -0.92151002, -0.49282334]])
```

In [37]:

```
1 # Get our predicted insurance rate for our new customer
2 random_forest_reg.predict(input_data)
```

Out[37]:

```
array([5989.69398529])
```

In [ ]:

```
1
```