

In [1]:

```

1 #Importing dataset
2
3 import pandas as pd
4 hr_df = pd.read_csv("D:\Data Science and Deep Learning for Business\datascienceforbu
5
6 print(hr_df.info())
7 print("\n",hr_df.shape)

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   employee_id           14999 non-null  int64
1   number_project        14999 non-null  int64
2   average_monthly_hours 14999 non-null  int64
3   time_spend_company    14999 non-null  int64
4   work_accident         14999 non-null  int64
5   left                 14999 non-null  int64
6   promotion_last_5years 14999 non-null  int64
7   department            14999 non-null  object
8   salary               14999 non-null  object
dtypes: int64(7), object(2)
memory usage: 1.0+ MB
None

```

(14999, 9)

In [2]:

```

1 #view categorical data and viewing their unique value
2
3 print(hr_df.select_dtypes(exclude=['int', 'float']).columns)
4 print("\n",hr_df['department'].unique())
5 print("\n",hr_df['salary'].unique())

```

Index(['department', 'salary'], dtype='object')

```

['sales' 'accounting' 'hr' 'technical' 'support' 'management' 'IT'
'product_mng' 'marketing' 'RandD']

```

```

['low' 'medium' 'high']

```

In [3]:

```
1 #Loading employee satisfaction dataset
2
3 emp_stats = pd.read_excel("D:\Data Science and Deep Learning for Business\datascience\employee_satisfaction.xlsx")
4 emp_stats.head()
```

Out[3]:

	EMPLOYEE #	satisfaction_level	last_evaluation
0	1003	0.38	0.53
1	1005	0.80	0.86
2	1486	0.11	0.88
3	1038	0.72	0.87
4	1057	0.37	0.52

In [4]:

```
1 #joining two datasets based on employee id column
2 main_df = hr_df.set_index('employee_id').join(emp_stats.set_index('EMPLOYEE #'))
3 main_df.head()
```

Out[4]:

	employee_id	number_project	average_monthly_hours	time_spend_company	Work_accident	last_evaluation
0	1003	2	157	3	0	0.53
1	1005	5	262	6	0	0.86
2	1486	7	272	4	0	0.88
3	1038	5	223	5	0	0.87
4	1057	2	159	3	0	0.52

In [5]:

```
1 #flattening the dataset after merge
2 main_df = main_df.reset_index()
3 main_df.head()
```

Out[5]:

	employee_id	number_project	average_monthly_hours	time_spend_company	Work_accident
0	1003	2	157	3	C
1	1005	5	262	6	C
2	1486	7	272	4	C
3	1038	5	223	5	C
4	1057	2	159	3	C



In [6]:

```
1 #checking any NULL values or missing values
2 print(main_df.isnull().sum())
3 #viewing null value records
4 main_df[main_df.isnull().any(axis=1)]
5
6
```

```
employee_id      0
number_project    0
average_monthly_hours  0
time_spend_company  0
work_accident     0
left              0
promotion_last_5years  0
department        0
salary            0
satisfaction_level 27
last_evaluation    27
dtype: int64
```

Out[6]:

	employee_id	number_project	average_monthly_hours	time_spend_company	Work_acc
	18	3794	2	160	3
	19	1140	5	262	5
	33	1230	2	140	3
	53	1340	2	132	3
	72	22316	2	149	3
	92	1581	2	143	3
	107	17376	2	148	3
	120	1739	4	158	4
	137	1847	2	129	3
	175	32923	4	164	2
	191	2160	4	226	6
	352	3150	4	262	6
	376	3250	4	296	2
	402	3405	5	275	5
	427	78130	3	180	4
	442	3635	5	229	5
	468	3755	5	245	5
	543	4150	5	237	5
	892	43615	4	276	5
	1588	42185	5	264	5

In [7]:

```
1 #filling missing value with mean values
2 main_df.fillna(main_df.mean(numeric_only=True),inplace=True)
3 main_df.head()
4
```

Out[7]:

	3609	21580	3	263	2
	employee_id	number_project	average_monthly_hours	time_spend_company	Work_accident
0	4122	100324505	23	157192	33C
1	4740	100527950	53	262253	63C
2	5028	148629640	74	272180	44C
3	6453	103838090	55	223166	52C
4	7005	105741535	24	159150	33C
	8630	50960	4	167	3
	9455	55770	4	270	3
	9901	58630	5	252	3
	10647	62595	4	165	3
	10962	64350	5	233	3

In [8]:

```

11575      20215      3      192      7
1 # Removing employee ID feature
11967 main_df_final = main_df.drop(columns='employee_id')
3 #print(main_df_final.head())
12422      72475      5      257      5
4 #viewing target column counts
12853 print(main_df_final['left'].value_counts())
6 # 1 - employee left; 0 - employee not left
13482      78780      3      207      7

01392511428  81315      3      133      3
1      3571
Name: left, dtype: int64

```

In [9]:

```

1 # Perform One Hot Encoding on Categorical Data
2
3 categorial = ['department', 'salary']
4 main_df_final = pd.get_dummies(main_df_final, columns=categorial, drop_first=True)
5 main_df_final.head()
6

```

Out[9]:

	number_project	average_monthly_hours	time_spend_company	Work_accident	left	promot
0	2	157	3	0	1	
1	5	262	6	0	1	
2	7	272	4	0	1	
3	5	223	5	0	1	
4	2	159	3	0	1	

In [10]:

```

1 #splitting the data
2
3 from sklearn.model_selection import train_test_split
4
5 #Removing target column
6 X = main_df_final.drop(['left'],axis=1).values
7 # Assigning target column
8 y = main_df_final['left'].values
9 # Splitting the data
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

```

In [11]:

```

1 # Normalize the data
2 from sklearn.preprocessing import StandardScaler
3
4 sc = StandardScaler()
5 X_train = sc.fit_transform(X_train)
6 X_test = sc.transform(X_test)

```

In [12]:

```

1 #Logistic Regression
2
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
5
6 model = LogisticRegression()
7 model.fit(X_train, y_train)
8
9 predictions = model.predict(X_test)
10
11 print("Accuracy {:.2f}%".format(100*accuracy_score(predictions, y_test)))
12 print("\nConfusion Matrix\n", confusion_matrix(y_test, predictions))
13 print(classification_report(y_test, predictions))
14
15

```

Accuracy 78.47%

Confusion Matrix

```

[[3142  249]
 [ 720  389]]

```

	precision	recall	f1-score	support
0	0.81	0.93	0.87	3391
1	0.61	0.35	0.45	1109
accuracy			0.78	4500
macro avg	0.71	0.64	0.66	4500
weighted avg	0.76	0.78	0.76	4500

In [13]:

```

1  #Deep Learning
2  import tensorflow.keras
3  from tensorflow.keras.models import Sequential
4  from tensorflow.keras.layers import Dense
5  from tensorflow.keras.regularizers import l2
6  from tensorflow.keras.layers import Dropout
7
8  model2 = Sequential()
9
10 # Hidden Layer 1
11 model2.add(Dense(270, activation='relu', input_dim=18, kernel_regularizer=l2(0.01)))
12 model2.add(Dropout(0.3))
13
14 # Hidden Layer 1
15 model2.add(Dense(180, activation='relu', input_dim=18, kernel_regularizer=l2(0.01)))
16 model2.add(Dropout(0.3))
17
18 # Hidden Layer 2
19 model2.add(Dense(90, activation = 'relu', input_dim=18, kernel_regularizer=l2(0.01)))
20 model2.add(Dropout(0.3))
21
22
23 model2.add(Dense(1, activation='sigmoid'))
24
25 model2.summary()
26
27 model2.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 270)	5130
dropout (Dropout)	(None, 270)	0
dense_1 (Dense)	(None, 180)	48780
dropout_1 (Dropout)	(None, 180)	0
dense_2 (Dense)	(None, 90)	16290
dropout_2 (Dropout)	(None, 90)	0
dense_3 (Dense)	(None, 1)	91
=====		
Total params: 70,291		
Trainable params: 70,291		
Non-trainable params: 0		
=====		



In [14]:

```
1 # training
2 batch_size = 10
3 epochs = 25
4
5 history = model2.fit(X_train,
6                     y_train,
7                     batch_size = batch_size,
8                     epochs = epochs,
9                     verbose = 1,
10                    validation_data = (X_test, y_test))
11
12 score = model2.evaluate(X_test, y_test, verbose=0)
13 print('Test loss:', score[0])
14 print('Test accuracy:', score[1])
```

```
Epoch 1/25
1050/1050 [=====] - 4s 3ms/step - loss: 0.6867 -
accuracy: 0.9007 - val_loss: 0.2846 - val_accuracy: 0.9498
Epoch 2/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2869 -
accuracy: 0.9377 - val_loss: 0.2444 - val_accuracy: 0.9507
Epoch 3/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2734 -
accuracy: 0.9380 - val_loss: 0.2374 - val_accuracy: 0.9524
Epoch 4/25
1050/1050 [=====] - 5s 5ms/step - loss: 0.2638 -
accuracy: 0.9403 - val_loss: 0.2247 - val_accuracy: 0.9593
Epoch 5/25
1050/1050 [=====] - 5s 5ms/step - loss: 0.2578 -
accuracy: 0.9412 - val_loss: 0.2338 - val_accuracy: 0.9476
Epoch 6/25
1050/1050 [=====] - 5s 5ms/step - loss: 0.2582 -
accuracy: 0.9409 - val_loss: 0.2297 - val_accuracy: 0.9507
Epoch 7/25
1050/1050 [=====] - 5s 5ms/step - loss: 0.2546 -
accuracy: 0.9412 - val_loss: 0.2206 - val_accuracy: 0.9576
Epoch 8/25
1050/1050 [=====] - 5s 5ms/step - loss: 0.2491 -
accuracy: 0.9412 - val_loss: 0.2176 - val_accuracy: 0.9544
Epoch 9/25
1050/1050 [=====] - 5s 4ms/step - loss: 0.2492 -
accuracy: 0.9445 - val_loss: 0.2182 - val_accuracy: 0.9587
Epoch 10/25
1050/1050 [=====] - 5s 4ms/step - loss: 0.2430 -
accuracy: 0.9433 - val_loss: 0.2157 - val_accuracy: 0.9529
Epoch 11/25
1050/1050 [=====] - 5s 4ms/step - loss: 0.2433 -
accuracy: 0.9445 - val_loss: 0.2073 - val_accuracy: 0.9580
Epoch 12/25
1050/1050 [=====] - 5s 4ms/step - loss: 0.2406 -
accuracy: 0.9446 - val_loss: 0.2059 - val_accuracy: 0.9567
Epoch 13/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2407 -
accuracy: 0.9429 - val_loss: 0.2098 - val_accuracy: 0.9531
Epoch 14/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2388 -
accuracy: 0.9436 - val_loss: 0.2177 - val_accuracy: 0.9518
Epoch 15/25
1050/1050 [=====] - 5s 4ms/step - loss: 0.2386 -
accuracy: 0.9454 - val_loss: 0.2006 - val_accuracy: 0.9589
Epoch 16/25
1050/1050 [=====] - 5s 5ms/step - loss: 0.2371 -
accuracy: 0.9436 - val_loss: 0.2088 - val_accuracy: 0.9589
Epoch 17/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2397 -
accuracy: 0.9444 - val_loss: 0.2016 - val_accuracy: 0.9587
Epoch 18/25
1050/1050 [=====] - 5s 5ms/step - loss: 0.2384 -
accuracy: 0.9436 - val_loss: 0.2034 - val_accuracy: 0.9584
Epoch 19/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2342 -
accuracy: 0.9459 - val_loss: 0.1956 - val_accuracy: 0.9596
Epoch 20/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2315 -
accuracy: 0.9480 - val_loss: 0.2192 - val_accuracy: 0.9529
Epoch 21/25
```

```

1050/1050 [=====] - 4s 4ms/step - loss: 0.2353 -
accuracy: 0.9445 - val_loss: 0.1934 - val_accuracy: 0.9631
Epoch 22/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2310 -
accuracy: 0.9469 - val_loss: 0.1924 - val_accuracy: 0.9622
Epoch 23/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2338 -
accuracy: 0.9453 - val_loss: 0.1953 - val_accuracy: 0.9618
Epoch 24/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2318 -
accuracy: 0.9457 - val_loss: 0.1960 - val_accuracy: 0.9598
Epoch 25/25
1050/1050 [=====] - 4s 4ms/step - loss: 0.2286 -
accuracy: 0.9470 - val_loss: 0.2119 - val_accuracy: 0.9511
Test loss: 0.2119394838809967
Test accuracy: 0.9511111378669739

```

In [15]:

```

1 predictions = model2.predict(X_test)
2 predictions = (predictions > 0.5)
3
4 print(confusion_matrix(y_test, predictions))
5 print(classification_report(y_test, predictions))

```

```


141/141 [=====] - 0s 2ms/step
[[3334   57]
 [ 163  946]]

```

	precision	recall	f1-score	support
0	0.95	0.98	0.97	3391
1	0.94	0.85	0.90	1109
accuracy			0.95	4500
macro avg	0.95	0.92	0.93	4500
weighted avg	0.95	0.95	0.95	4500


In [16]:

```
1 # Loss charts
2 import matplotlib.pyplot as plt
3
4 history_dict = history.history
5
6 loss_values = history_dict['loss']
7 val_loss_values = history_dict['val_loss']
8 epochs = range(1, len(loss_values) + 1)
9
10 line1 = plt.plot(epochs, val_loss_values, label='Validation/Test Loss')
11 line2 = plt.plot(epochs, loss_values, label='Training Loss')
12 plt.setp(line1, linewidth=2.0, marker = '+', markersize=10.0)
13 plt.setp(line2, linewidth=2.0, marker = 'o', markersize=10.0)
14 plt.xlabel('Epochs')
15 plt.ylabel('Loss')
16 plt.grid(True)
17 plt.legend()
18 plt.show()
```



In [17]:

```
1 # accuracy charts
2 import matplotlib.pyplot as plt
3
4 history_dict = history.history
5
6 acc_values = history_dict['accuracy']
7 val_acc_values = history_dict['val_accuracy']
8 epochs = range(1, len(loss_values) + 1)
9
10 line1 = plt.plot(epochs, val_acc_values, label='Validation/Test Accuracy')
11 line2 = plt.plot(epochs, acc_values, label='Training Accuracy')
12 plt.setp(line1, linewidth=2.0, marker = '+', markersize=10.0)
13 plt.setp(line2, linewidth=2.0, marker = 'o', markersize=10.0)
14 plt.xlabel('Epochs')
15 plt.ylabel('Accuracy')
16 plt.grid(True)
17 plt.legend()
18 plt.show()
```



In [ ]:

1